

```
In [1]: import numpy as np
        from sklearn.metrics import confusion_matrix

In [2]: # Load data set and code labels as 0 = 'NO', 1 = 'DH', 2 = 'SL'
        labels = [b'NO', b'DH', b'SL']
        data = np.loadtxt('spine-data.txt', converters = {6: lambda s: labels.index(s)})

In [3]: # Separate features from labels
        X = data[:,0:6]
        y = data[:,6]

        # Split the data into a training set, consisting of the first 250 points,
        # and a test set, consisting of the remaining 60 points.
        train_data = X[list(range(0,250)),:]
        train_label = y[list(range(0,250))]
        test_data = X[list(range(250,310)),:]
        test_label = y[list(range(250,310))]

In [4]: def distance_L1(x, y):
        return np.sum(np.abs(x-y))

        def square_distance_L2(x, y):
            return np.sum(np.square(x-y))

        def predict(x, y):
            return y[np.argmin(x)]

        def NN_Classifier_L1(trainx, trainy, testx):
            testy_L1 = []
            for i in range(len(testx)):
                distance = [distance_L1(testx[i], trainx[j]) for j in range(len(trainx))]
                test_predicted = predict(distance, trainy)
                testy_L1.append(test_predicted)
            return np.asarray(testy_L1)

        def NN_Classifier_L2(trainx, trainy, testx):
            testy_L2 = []
            for i in range(len(testx)):
                distance = [square_distance_L2(testx[i], trainx[j]) for j in range(len(trainx))]
                test_predicted = predict(distance, trainy)
                testy_L2.append(test_predicted)
            return np.asarray(testy_L2)

        def error_rate(testy, testy_fit):
            return float(sum(testy != testy_fit))/len(testy)
```

(a) What error rates do you get on the test set for each of the two distance functions?

```
In [5]: test_label_L1 = NN_Classifier_L1(train_data, train_label, test_data)
        test_label_L2 = NN_Classifier_L2(train_data, train_label, test_data)

        error_L1 = error_rate(test_label, test_label_L1)
        error_L2 = error_rate(test_label, test_label_L2)

        print("Error rate of NN_Classifier_L1:\n", error_L1)
        print("\nError rate of NN_Classifier_L2:\n", error_L2)

Error rate of NN_Classifier_L1:
0.21666666666666667

Error rate of NN_Classifier_L2:
0.23333333333333334
```

(b) For each of the two distance functions, give the confusion matrix of the NN classifier.

```
In [6]: print('Confusion matrix of nearest neighbor classifier L1:\n', confusion_matrix(test_label, test_label_L1))

Confusion matrix of nearest neighbor classifier L1:
[[14  0  2]
 [ 9  9  0]
 [ 1  1 24]]

In [7]: print('Confusion matrix of nearest neighbor classifier L2:\n', confusion_matrix(test_label, test_label_L2))

Confusion matrix of nearest neighbor classifier L2:
[[12  1  3]
 [ 9  9  0]
 [ 1  0 25]]

In [ ]:
```

